

In this notebook, You will do amazon review classification with BERT.[Download data from [this link](#)]

It contains 5 parts as below. Detailed instructions are given in the each cell. please read every comment we have written.

1. Preprocessing
2. Creating a BERT model from the Tensorflow HUB.
3. Tokenization
4. getting the pretrained embedding Vector for a given review from the BERT.
5. Using the embedding data apply NN and classify the reviews.
6. Creating a Data pipeline for BERT Model.

instructions:

1. Don't change any Grader Functions. Don't manipulate any Grader functions. If you manipulate any, it will be considered as plagiarised.
2. Please read the instructions on the code cells and markdown cells. We will explain what to write.
3. please return outputs in the same format what we asked. Eg. Don't return List if we are asking for a numpy array.
4. Please read the external links that we are given so that you will learn the concept behind the code that you are writing.
5. We are giving instructions at each section if necessary, please follow them.

Every Grader function has to return True.

```
# Downloading dataset zip file
```

```
!gdown --id 1ziSQS7bTrc0mV1I1YbzNpYrL452VVol-  
Downloading...  
From: https://drive.google.com/uc?id=1ziSQS7bTrc0mV1I1YbzNpYrL452VVol-  
To: /content/Reviews.csv.zip  
120MB [00:01, 108MB/s]
```

```
# Unzipping file
```

```
!unzip 'Reviews.csv.zip'  
Archive:  Reviews.csv.zip  
  inflating: Reviews.csv
```

```
#all imports  
import numpy as np  
import pandas as pd  
import tensorflow as tf  
import tensorflow_hub as hub  
from tensorflow.keras.models import Model  
from tqdm import tqdm  
import re  
from sklearn.metrics import roc_auc_score  
from sklearn.model_selection import train_test_split  
import matplotlib.pyplot as plt
```

```
tf.test.gpu_device_name()
```

```
 '/device:GPU:0'
```

Grader function 1

```
def grader_tf_version():  
    assert((tf.__version__)>'2')  
    return True  
grader_tf_version()
```

True

In []:

In []:

In []:

In []:

Out[]:

In []:

Out[]:

Part-1: Preprocessing

In []:

```
#Read the dataset - Amazon fine food reviews
reviews = pd.read_csv("Reviews.csv")
#check the info of the dataset
reviews.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 568454 entries, 0 to 568453
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                     568454 non-null  int64
1   ProductId              568454 non-null  object
2   UserId                 568454 non-null  object
3   ProfileName            568438 non-null  object
4   HelpfulnessNumerator    568454 non-null  int64
5   HelpfulnessDenominator  568454 non-null  int64
6   Score                  568454 non-null  int64
7   Time                   568454 non-null  int64
8   Summary                568427 non-null  object
9   Text                   568454 non-null  object
dtypes: int64(5), object(5)
memory usage: 43.4+ MB
```

In []:

```
#get only 2 columns - Text, Score
#drop the NAN values
reviews = reviews[['Text', 'Score']]
reviews.dropna(inplace=True)
```

In []:

```
#if score > 3, set score = 1
#if score <= 2, set score = 0
#if score == 3, remove the rows.
reviews = reviews[reviews.Score != 3]
def partition(x):
    if x < 3:
        return 0
    return 1

scores = reviews['Score']
newScores = scores.map(partition)
reviews['Score'] = newScores
```

Grader function 2

In []:

```
def grader_reviews():
    temp_shape = (reviews.shape == (525814, 2)) and (reviews.Score.value_counts()[1] == 443777)
    assert(temp_shape == True)
    return True
grader_reviews()
```

True

Out []:

In []:

```
# Adding column of word-length

def get_wordlen(x):
    return len(x.split())
reviews['len'] = reviews.Text.apply(get_wordlen)
reviews = reviews[reviews.len < 50]
reviews = reviews.sample(n=100000, random_state=30)
```

In []:

```
#remove HTML from the Text column and save in the Text column only
def clean_text(x):
    return re.sub(r"http\S+", "", x)

reviews['Text'] = reviews.Text.apply(clean_text)
```

In []:

```
#print head 5
reviews.head()
```

Out[]:

	Text	Score	len
64117	The tea was of great quality and it tasted lik...	1	30
418112	My cat loves this. The pellets are nice and s...	1	31
357829	Great product. Does not completely get rid of ...	1	41
175872	This gum is my favorite! I would advise every...	1	27
178716	I also found out about this product because of...	1	22

In []:

```
#split the data into train and test data(20%) with Stratify sampling, random state 33,
X_train, X_test, y_train, y_test = train_test_split(reviews[['Text', 'len']], reviews['Score'], test_size

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(80000, 2)
(20000, 2)
(80000,)
(20000,)
```

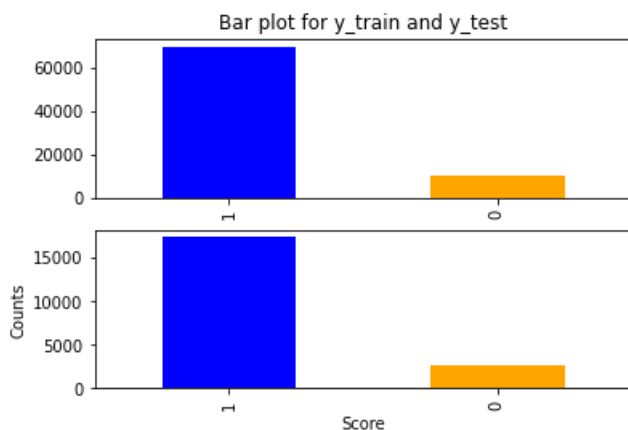
In []:

```
#plot bar graphs of y_train and y_test
plt.figure(1)

plt.subplot(211)
plt.title("Bar plot for y_train and y_test")
y_train.value_counts().plot(kind='bar', color=('blue','orange'))
plt.subplot(212)
y_test.value_counts().plot(kind='bar', color=('blue','orange'))

plt.xlabel("Score")
plt.ylabel("Counts")

plt.show()
```



We have more positive reviews than negative reviews.

In []:

```
#saving to disk. if we need, we can load preprocessed data directly.
reviews.to_csv('preprocessed.csv', index=False)
```

Part-2: Creating BERT Model

If you want to know more about BERT, You can watch live sessions on Transformers and BERT. we will strongly recommend you to read [Transformers](#), [BERT Paper](#) and, [This blog](#).

For this assignment, we are using [BERT uncased Base model](#). It uses L=12 hidden layers (i.e., Transformer blocks), a hidden size of H=768, and A=12 attention heads.

In []:

```
## Loading the Pretrained Model from tensorflow HUB
tf.keras.backend.clear_session()
```

```
# maximum length of a seq in the data we have, for now i am making it as 55. You can change this
max_seq_length = 55

#BERT takes 3 inputs

#this is input words. Sequence of words represented as integers
input_word_ids = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_word_ids")

#mask vector if you are padding anything
input_mask = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_mask")

#segment vectors. If you are giving only one sentence for the classification, total seg vector is 0.
#If you are giving two sentenced with [sep] token separated, first seq segment vectors are zeros and
#second seq segment vector are 1's
segment_ids = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="segment_ids")

#bert layer
bert_layer = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/1", trainable=False)
pooled_output, sequence_output = bert_layer([input_word_ids, input_mask, segment_ids])

#Bert model
#We are using only pooled output not sequence out.
#If you want to know about those, please read https://www.kaggle.com/questions-and-answers/86510
bert_model = Model(inputs=[input_word_ids, input_mask, segment_ids], outputs=pooled_output)
```

In []:

```
bert_model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_word_ids (InputLayer)	[(None, 55)]	0	
input_mask (InputLayer)	[(None, 55)]	0	
segment_ids (InputLayer)	[(None, 55)]	0	
keras_layer (KerasLayer)	[(None, 768), (None, 109482241		input_word_ids[0][0] input_mask[0][0] segment_ids[0][0]

=====
Total params: 109,482,241
Trainable params: 0
Non-trainable params: 109,482,241

In []:

```
bert_model.output
```

Out []:

```
<KerasTensor: shape=(None, 768) dtype=float32 (created by layer 'keras_layer')>
```

Part-3: Tokenization

In []:

```
#getting Vocab file
vocab_file = bert_layer.resolved_object.vocab_file.asset_path.numpy()
do_lower_case = bert_layer.resolved_object.do_lower_case.numpy()
```

In []:

```
#import tokenization - We have given tokenization.py file
!pip install sentencepiece
from tokenization import FullTokenizer
```

```
Collecting sentencepiece
  Downloading
https://files.pythonhosted.org/packages/14/67/e42bd1181472c95c8cda79305df848264f2a7f62740995a46945d9797b6
tencepiece-0.1.95-cp36-cp36m-manylinux2014_x86_64.whl (1.2MB)
|████████████████████████████████████████| 1.2MB 5.4MB/s
Installing collected packages: sentencepiece
Successfully installed sentencepiece-0.1.95
```

In []:

```
# Create tokenizer " Instantiate FullTokenizer"
# name must be "tokenizer"
# the FullTokenizer takes two parameters 1. vocab_file and 2. do_lower_case
# we have created these in the above cell ex: FullTokenizer(vocab_file, do_lower_case )
```

```
# please check the "tokenization.py" file the complete implementation
```

```
tokenizer = FullTokenizer(vocab_file, do_lower_case)
```

Grader function 3

In []:

```
#it has to give no error
def grader_tokenize(tokenizer):
    out = False
    try:
        out=('[CLS]' in tokenizer.vocab) and ('[SEP]' in tokenizer.vocab)
    except:
        out = False
    assert(out==True)
    return out
grader_tokenize(tokenizer)
```

True

Out[]:

In []:

```
# Create train and test tokens (X_train_tokens, X_test_tokens) from (X_train, X_test) using Tokenizer and
# add '[CLS]' at start of the Tokens and '[SEP]' at the end of the tokens.
# maximum number of tokens is 55(We already given this to BERT layer above) so shape is (None, 55)
# if it is less than 55, add '[PAD]' token else truncate the tokens length.(similar to padding)

# Based on padding, create the mask for Train and Test ( 1 for real token, 0 for '[PAD]'),
# it will also same shape as input tokens (None, 55) save those in X_train_mask, X_test_mask

# Create a segment input for train and test. We are using only one sentence so all zeros. This shape wil.

# type of all the above arrays should be numpy arrays

# after execution of this cell, you have to get
# X_train_tokens, X_train_mask, X_train_segment
# X_test_tokens, X_test_mask, X_test_segment

def tokenize(texts):
    tokens = tokenizer.tokenize(text)
    if (len(tokens) < max_seq_length-2):
        tokens.extend(['[PAD]'] * (max_seq_length-2 - len(tokens)))
    elif (len(tokens) > max_seq_length-2):
        tokens = tokens[:max_seq_length-2]
    tokens = ['[CLS]', *tokens, '[SEP]']
    token_ids = np.array(tokenizer.convert_tokens_to_ids(tokens))
    mask = np.array([0 if i=='[PAD]' else 1 for i in tokens])
    segment = np.array([0]*max_seq_length)
    return token_ids, mask, segment

X_train_tokens = []
X_train_mask = []
X_train_segment = []
for text in tqdm(X_train['Text'].values):
    tokens, mask, segment = tokenize(text)
    X_train_tokens.append(tokens)
    X_train_mask.append(mask)
    X_train_segment.append(segment)

X_test_tokens = []
X_test_mask = []
X_test_segment = []
for text in tqdm(X_test['Text'].values):
    tokens, mask, segment = tokenize(text)
    X_test_tokens.append(tokens)
    X_test_mask.append(mask)
    X_test_segment.append(segment)

X_train_tokens = np.array(X_train_tokens)
X_train_mask = np.array(X_train_mask)
X_train_segment = np.array(X_train_segment)
X_test_tokens = np.array(X_test_tokens)
X_test_mask = np.array(X_test_mask)
X_test_segment = np.array(X_test_segment)
```



```

out = False

if type(X_train_tokens) == np.ndarray:

    temp_shapes = (X_train_tokens.shape[1]==max_seq_length) and (X_train_mask.shape[1]==max_seq_length)
    (X_train_segment.shape[1]==max_seq_length)

    segment_temp = not np.any(X_train_segment)

    mask_temp = np.sum(X_train_mask==0) == np.sum(X_train_tokens==0)

    no_cls = np.sum(X_train_tokens==tokenizer.vocab['[CLS]'])==X_train_tokens.shape[0]

    no_sep = np.sum(X_train_tokens==tokenizer.vocab['[SEP]'])==X_train_tokens.shape[0]

    out = temp_shapes and segment_temp and mask_temp and no_cls and no_sep

else:
    print('Type of all above token arrays should be numpy array not list')
    out = False
assert(out==True)
return out

grader_alltokens_train()

```

True

Out[]:

Grader function 5

In []:

```

def grader_alltokens_test():
    out = False
    if type(X_test_tokens) == np.ndarray:

        temp_shapes = (X_test_tokens.shape[1]==max_seq_length) and (X_test_mask.shape[1]==max_seq_length)
        (X_test_segment.shape[1]==max_seq_length)

        segment_temp = not np.any(X_test_segment)

        mask_temp = np.sum(X_test_mask==0) == np.sum(X_test_tokens==0)

        no_cls = np.sum(X_test_tokens==tokenizer.vocab['[CLS]'])==X_test_tokens.shape[0]

        no_sep = np.sum(X_test_tokens==tokenizer.vocab['[SEP]'])==X_test_tokens.shape[0]

        out = temp_shapes and segment_temp and mask_temp and no_cls and no_sep

    else:
        print('Type of all above token arrays should be numpy array not list')
        out = False
    assert(out==True)
    return out
grader_alltokens_test()

```

True

Out[]:

Part-4: Getting Embeddings from BERT Model

We already created the BERT model in the part-2 and input data in the part-3. We will utilize those two and will get the embeddings for each sentence in the Train and test data.

In []:

```
bert_model.input
```

Out[]:

```

<KerasTensor: shape=(None, 55) dtype=int32 (created by layer 'input_word_ids')>,
<KerasTensor: shape=(None, 55) dtype=int32 (created by layer 'input_mask')>,
<KerasTensor: shape=(None, 55) dtype=int32 (created by layer 'segment_ids')>]

```

In []:

```
bert_model.output
```

Out[]:

```
<KerasTensor: shape=(None, 768) dtype=float32 (created by layer 'keras_layer')>
```

In []:

```
# get the train output, BERT model will give one output so save in
```

```
# X_train_pooled_output
X_train_pooled_output=bert_model.predict([X_train_tokens,X_train_mask,X_train_segment])

# get the test output, BERT model will give one output so save in
# X_test_pooled_output
X_test_pooled_output=bert_model.predict([X_test_tokens,X_test_mask,X_test_segment])

##save all your results to disk so that, no need to run all again.
pickle.dump((X_train_pooled_output, X_test_pooled_output),open('final_output.pkl','wb'))

#X_train_pooled_output, X_test_pooled_output= pickle.load(open('final_output.pkl', 'rb'))
```

In []:

In []:

In []:

Grader function 6

In []:

```
#now we have X_train_pooled_output, y_train
#X_test_pooled_output, y_test

#please use this grader to evaluate
def greader_output():
    assert(X_train_pooled_output.shape[1]==768)
    assert(len(y_train)==len(X_train_pooled_output))
    assert(X_test_pooled_output.shape[1]==768)
    assert(len(y_test)==len(X_test_pooled_output))
    assert(len(y_train.shape)==1)
    assert(len(X_train_pooled_output.shape)==2)
    assert(len(y_test.shape)==1)
    assert(len(X_test_pooled_output.shape)==2)
    return True
grader_output()
```

True

Out []:

Part-5: Training a NN with 768 features

Create a NN and train the NN.

1. You have to use AUC as metric.
2. You can use any architecture you want.
3. You have to use tensorboard to log all your metrics and Losses. You have to send those logs.
4. Print the loss and metric at every epoch.
5. You have to submit without overfitting and underfitting.

In []:

```
##imports
from tensorflow.keras.layers import Input, Dense, Activation, Dropout, BatchNormalization
from tensorflow.keras.models import Model
from tensorflow.keras.regularizers import l2
from tensorflow.keras.callbacks import TensorBoard, Callback
import datetime
import os
```

In []:

```
##create a NN on top of BERT model (It will take the BERT model output as its input)
# 7 dense layers are used with final output layer to create the NN. Each layer uses regularizer to
# Dropout and Batch Normalization are used after few layers to avoid overfitting and normalize values re.
# As its a deep NN, Batch Normalization helps to get better result
# As its a binary classification, 'sigmoid' activation function is used at output layer
# In other dense layers, 'relu' activation is used with 'he_normal' kernel initializer.
```

```
BERT_output = Input(shape=(768,), name="BERT_output")
```

```
dense1 = Dense(1024, activation='relu', kernel_initializer='he_normal', kernel_regularizer=l2(0.0001), name='dense1')
dense2 = Dense(512, activation='relu', kernel_initializer='he_normal', kernel_regularizer=l2(0.0001), name='dense2')
dense3 = Dense(256, activation='relu', kernel_initializer='he_normal', kernel_regularizer=l2(0.0001), name='dense3')
dropout1 = Dropout(0.5)(dense3)
bn1 = BatchNormalization()(dropout1)
dense4 = Dense(128, activation='relu', kernel_initializer='he_normal', kernel_regularizer=l2(0.0001), name='dense4')
dense5 = Dense(64, activation='relu', kernel_initializer='he_normal', kernel_regularizer=l2(0.0001), name='dense5')
dense6 = Dense(32, activation='relu', kernel_initializer='he_normal', kernel_regularizer=l2(0.0001), name='dense6')
dense7 = Dense(16, activation='relu', kernel_initializer='he_normal', kernel_regularizer=l2(0.0001), name='dense7')
dropout2 = Dropout(0.5)(dense7)
bn2 = BatchNormalization()(dropout2)
output = Dense(1, activation='sigmoid', kernel_regularizer=l2(0.0001), name='output')(bn2)
```



```

model = Model(inputs=BERT_output, outputs=output)
model.summary()

```

```

Model: "model_1"

```

Layer (type)	Output Shape	Param #
BERT_output (InputLayer)	[(None, 768)]	0
dense1 (Dense)	(None, 1024)	787456
dense2 (Dense)	(None, 512)	524800
dense3 (Dense)	(None, 256)	131328
dropout (Dropout)	(None, 256)	0
batch_normalization (BatchNo	(None, 256)	1024
dense4 (Dense)	(None, 128)	32896
dense5 (Dense)	(None, 64)	8256
dense6 (Dense)	(None, 32)	2080
dense7 (Dense)	(None, 16)	528
dropout_1 (Dropout)	(None, 16)	0
batch_normalization_1 (Batch	(None, 16)	64
output (Dense)	(None, 1)	17
Total params: 1,488,449		
Trainable params: 1,487,905		
Non-trainable params: 544		

```

In []:

```

```

# auc function to be passed as metric

def auc(y_true, y_pred):
    try:
        return tf.py_function(roc_auc_score, (y_true, y_pred), tf.double)
    except ValueError:
        return 0.5

```

```

In []:

```

```

# Custom callback to be used for manipulating the trainig life-cycle

class My_callback(Callback):
    '''Custom callback to save best model, stop training if accuracy is not improving and compute F1-sco:
    def __init__(self, epochs):
        self.epochs = epochs

    def save_best(self):
        # Saving the best model based on high val_accuracy
        best_val_auc = max(self.history['model'].keys())
        print("Saving best weights before terminating having val_auc {}".format(best_val_auc))
        best_model = self.history['model'].get(best_val_auc)
        filepath="model.hdf5"
        best_model.save(filepath)

    def on_train_begin(self, logs={}):
        # On begin of training, history dict is created with keys [model, loss, accuracy, val_loss, val_
        self.history={'model': dict(), 'val_auc': []}

    def on_epoch_end(self, epoch, logs={}):
        # In each epoch end, update loss, accuracy and F1-score
        loss = logs.get('loss')
        if (loss is not None):
            if (np.isnan(loss) or np.isinf(loss)):
                print('Invalid loss. Terminating at epoch {}'.format(epoch))
            else:
                # Record model at each epcoh with val_auc
                if logs.get('val_auc', -1) != -1:
                    self.history['val_auc'].append(logs.get('val_auc'))
                    self.history['model'][logs.get('val_auc')] = self.model

```

```

# Stop training if val_auc doesnt improve for last 4 epochs
if len(self.history['val_auc']) >= 4:
    if ((self.history['val_auc'][-1] < self.history['val_auc'][-4]) and
        (self.history['val_auc'][-1] < self.history['val_auc'][-3]) and
        (self.history['val_auc'][-1] < self.history['val_auc'][-2])):
        print("Validation auc not improving and terminated at epoch {}".format(epoch))
        self.save_best()
        self.model.stop_training = True

# Finally save the best model at the end of training
if epoch == self.epochs-1:
    self.save_best()

```

In []:

```
%load_ext tensorboard
```

In []:

```

#Train NN with the BERT output values and Y-values (scores).
# Use 50 epochs, Adam optimizer with 0.001 learning rate, accuracy and auc as evaluation metrics.
# Use binary_crossentropy as loss function
# tensorboard and callback object needs to be passed as callback

epochs = 50
logdir = os.path.join("logs")
tensorboard_callback = TensorBoard(log_dir=logdir, histogram_freq=1, write_graph=True)
myCallback = My_callback(epochs)

opti = tf.keras.optimizers.Adam(learning_rate = 0.001)
model.compile(optimizer=opti, loss='binary_crossentropy', metrics=['accuracy', auc])
model.fit(X_train_pooled_output, y_train, epochs=epochs,
          validation_data=(X_test_pooled_output, y_test),
          batch_size=128, callbacks=[tensorboard_callback, myCallback])

```

```

Epoch 1/50
 3/625 [.....] - ETA: 28s - loss: 1.2871 - accuracy: 0.4727 - auc: 0.4867
WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time:
0.0053s vs `on_train_batch_end` time: 0.0132s). Check your callbacks.
WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time:
0.0053s vs `on_train_batch_end` time: 0.0132s). Check your callbacks.
625/625 [=====] - 5s 6ms/step - loss: 0.8903 - accuracy: 0.7044 - auc: 0.5012 -
val_loss: 0.5506 - val_accuracy: 0.8701 - val_auc: 0.5516
Epoch 2/50
625/625 [=====] - 3s 5ms/step - loss: 0.5309 - accuracy: 0.8703 - auc: 0.5101 -
val_loss: 0.4777 - val_accuracy: 0.8701 - val_auc: 0.9143
Epoch 3/50
625/625 [=====] - 3s 5ms/step - loss: 0.3715 - accuracy: 0.8806 - auc: 0.8720 -
val_loss: 0.2789 - val_accuracy: 0.8920 - val_auc: 0.9409
Epoch 4/50
625/625 [=====] - 3s 5ms/step - loss: 0.2970 - accuracy: 0.8991 - auc: 0.9100 -
val_loss: 0.3096 - val_accuracy: 0.8982 - val_auc: 0.9464
Epoch 5/50
625/625 [=====] - 3s 5ms/step - loss: 0.2611 - accuracy: 0.9077 - auc: 0.9260 -
val_loss: 0.8546 - val_accuracy: 0.5907 - val_auc: 0.9386
Epoch 6/50
625/625 [=====] - 4s 6ms/step - loss: 0.2408 - accuracy: 0.9111 - auc: 0.9342 -
val_loss: 0.2211 - val_accuracy: 0.9259 - val_auc: 0.9421
Epoch 7/50
625/625 [=====] - 3s 5ms/step - loss: 0.2307 - accuracy: 0.9132 - auc: 0.9344 -
val_loss: 0.3465 - val_accuracy: 0.8749 - val_auc: 0.9499
Epoch 8/50
625/625 [=====] - 3s 5ms/step - loss: 0.2223 - accuracy: 0.9146 - auc: 0.9381 -
val_loss: 0.2428 - val_accuracy: 0.8974 - val_auc: 0.9476
Epoch 9/50
625/625 [=====] - 3s 5ms/step - loss: 0.2175 - accuracy: 0.9166 - auc: 0.9397 -
val_loss: 0.1993 - val_accuracy: 0.9176 - val_auc: 0.9521
Epoch 10/50
625/625 [=====] - 3s 5ms/step - loss: 0.2174 - accuracy: 0.9151 - auc: 0.9398 -
val_loss: 0.5699 - val_accuracy: 0.8704 - val_auc: 0.9516
Epoch 11/50
625/625 [=====] - 3s 5ms/step - loss: 0.2127 - accuracy: 0.9161 - auc: 0.9381 -
val_loss: 0.3734 - val_accuracy: 0.8741 - val_auc: 0.9514
Epoch 12/50
625/625 [=====] - 3s 5ms/step - loss: 0.2067 - accuracy: 0.9196 - auc: 0.9431 -
val_loss: 0.3455 - val_accuracy: 0.8753 - val_auc: 0.9530
Epoch 13/50
625/625 [=====] - 3s 5ms/step - loss: 0.2072 - accuracy: 0.9193 - auc: 0.9418 -
val_loss: 0.1891 - val_accuracy: 0.9317 - val_auc: 0.9508
Validation auc not improving and terminated at epoch 13
Saving best weights before terminating having val_auc 0.9529764652252197

```

Out[]:


```
<tensorflow.python.keras.callbacks.History at 0x7f4d329ef3c8>
```

The model got trained and the best weight values are saved that gave 0.9529 val_auc. Now when we use this model after BERT, it can give better AUC score.

In []:


```
%tensorboard --logdir logs
```

epoch_accuracy



epoch_auc

epoch_auc



The training and validaion metrics got improved with each epoch and finally AUC score reached 0.95.

Part-6: Creating a Data pipeline for BERT Model

1. Download data from [here](#)
2. Read the csv file
3. Remove all the html tags
4. Now do tokenization [Part 3 as mentioned above]
 - Create tokens,mask array and segment array
5. Get Embeddings from BERT Model [Part 4 as mentioned above] , let it be X_test
 - Print the shape of output(X_test.shape).You should get (352,768)
6. Predict the output of X_test with the Neural network model which we trained earlier.
7. Print the occurrences of class labels in the predicted output

</pre>

In []:

```
# Building a pipeline where we can give the test data, new NN model and pre_trained BERT model as input.
# Threshold value will be passed to convert probability score to final output with high confidence.
# it will clean the input data and predict response with BERT model.
# The output will be fed into new NN model to get finally predicted output (Y_pred)
```

```
def predictBERT(data, model, bertModel, threshold):
    data['Text'] = data.Text.apply(clean_text)
    X = data['Text'].values
    X_tokens = []
    X_mask = []
    X_segment = []
    for text in tqdm(X):
        tokens, mask, segment = tokenize(text)
        X_tokens.append(tokens)
        X_mask.append(mask)
        X_segment.append(segment)
    X_tokens = np.array(X_tokens)
    X_mask = np.array(X_mask)
    X_segment = np.array(X_segment)
    bert_output = bertModel.predict([X_tokens, X_mask, X_segment])
    y_pred_prob = model.predict(bert_output)
    y_pred = []
    for y in y_pred_prob:
        if (y > threshold):
            y_pred.append(1)
        else:
            y_pred.append(0)
    return np.array(y_pred)
```

In []:

```
# Load test data from file
# Load the best weights we got during training to our NN
# Put a threshold to determine the final predicted review score
```

```
test_data = pd.read_csv("test.csv")
model.load_weights('model.hdf5')
threshold = 0.9
predictBERT(test_data, model, bert_model, threshold)
```

Out[]:

In [:

```
!zip -r logs.zip logs/

adding: logs/ (stored 0%)
adding: logs/validation/ (stored 0%)
adding: logs/validation/events.out.tfevents.1613972224.47d3b28fe05c.76.33831.v2 (deflated 59%)
adding: logs/train/ (stored 0%)
adding: logs/train/events.out.tfevents.1613972220.47d3b28fe05c.76.27825.v2 (deflated 86%)
adding: logs/train/events.out.tfevents.1613972221.47d3b28fe05c.profile-empty (deflated 5%)
adding: logs/train/plugins/ (stored 0%)
adding: logs/train/plugins/profile/ (stored 0%)
adding: logs/train/plugins/profile/2021_02_22_05_37_01/ (stored 0%)
adding: logs/train/plugins/profile/2021_02_22_05_37_01/47d3b28fe05c.overview_page.pb (deflated 57%)
adding: logs/train/plugins/profile/2021_02_22_05_37_01/47d3b28fe05c.tensorflow_stats.pb (deflated 74%)
adding: logs/train/plugins/profile/2021_02_22_05_37_01/47d3b28fe05c.trace.json.gz (deflated 1%)
adding: logs/train/plugins/profile/2021_02_22_05_37_01/47d3b28fe05c.memory_profile.json.gz (stored 0%)
adding: logs/train/plugins/profile/2021_02_22_05_37_01/47d3b28fe05c.input_pipeline.pb (deflated 55%)
adding: logs/train/plugins/profile/2021_02_22_05_37_01/47d3b28fe05c.xplane.pb (deflated 82%)
adding: logs/train/plugins/profile/2021_02_22_05_37_01/47d3b28fe05c.kernel_stats.pb (deflated 96%)
```