

```
# import libraries
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
IQ_train_df = pd.read_csv('train.csv')
IQ_train_df.head()
```

		Id	v2a1	hacdor	rooms	hacapo	v14a	refrig	v18q
v18q1 \									
0	ID_279628684	190000.0		0	3	0	1	1	0
NaN									
1	ID_f29eb3ddd	135000.0		0	4	0	1	1	1
1.0									
2	ID_68de51c94	NaN		0	8	0	1	1	0
NaN									
3	ID_d671db89c	180000.0		0	5	0	1	1	1
1.0									
4	ID_d56d6f5f5	180000.0		0	5	0	1	1	1
1.0									

	r4h1	...	SQBescolari	SQBage	SQBhogar_total	SQBedjefe
SQBhogar_nin \						
0	0	...	100	1849	1	100
0						
1	0	...	144	4489	1	144
0						
2	0	...	121	8464	1	0
0						
3	0	...	81	289	16	121
4						
4	0	...	121	1369	16	121
4						

	SQBovercrowding	SQBdependency	SQBmeaned	agesq	Target
0	1.000000	0.0	100.0	1849	4
1	1.000000	64.0	144.0	4489	4
2	0.250000	64.0	121.0	8464	4
3	1.777778	1.0	121.0	289	4
4	1.777778	1.0	121.0	1369	4

```
[5 rows x 143 columns]
```

```
IQ_train_df.shape
```

```
(9557, 143)
```

We have 9557 records with 143 columns.

### task 1: Identify the output variable

Target column is the output variable.

We can use `y = IQ_train_df['target']` later.

### Task 2: Understand the type of data.

```
IQ_train_df.dtypes.value_counts()
```

```
int64      130
float64      8
object       5
dtype: int64
```

There are 130 integer, 8 float and 5 object/string datatypes.

Lets find out the columns with datatype object.

```
IQ_train_df.columns[IQ_train_df.dtypes == 'object']
```

```
Index(['Id', 'idhogar', 'dependency', 'edjefe', 'edjefa'],
      dtype='object')
```

We can remove Id as it doesnt give any insights.

```
IQ_train_df.drop(['Id'], axis=1, inplace=True)
```

Lets check the object datatype data.

```
IQ_train_df[['idhogar', 'dependency', 'edjefe', 'edjefa']].head()
```

	idhogar	dependency	edjefe	edjefa
0	21eb7fcc1	no	10	no
1	0e5d7a658	8	12	no
2	2c7317ea8	8	no	11
3	2b58d945f	yes	11	no
4	2b58d945f	yes	11	no

```
IQ_train_df['idhogar'].value_counts()
```

```
fd8a6d014    13
0c7436de6    12
ae6cf0558    12
b7a0b59d7    11
3fe29a56b    11
```

```

3659d839d    1
d2e45f8ad    1
7c1cfa65c    1
3a3346eff    1
980b28caa    1
Name: idhogar, Length: 2988, dtype: int64

```

```

IQ_train_df[IQ_train_df['idhogar'] == 'fd8a6d014'].head()

```

	v2a1	hacdor	rooms	hacapo	v14a	refrig	v18q	v18q1	r4h1
r4h2 \									
2467	80000.0	1	4	1	1	1	0	NaN	5
2									
2468	80000.0	1	4	1	1	1	0	NaN	5
2									
2469	80000.0	1	4	1	1	1	0	NaN	5
2									
2470	80000.0	1	4	1	1	1	0	NaN	5
2									
2471	80000.0	1	4	1	1	1	0	NaN	5
2									

	... SQBescolari	SQBage	SQBhogar_total	SQBedjefe
SQBhogar_nin \				
2467 ...	49	256	169	0
81				
2468 ...	0	0	169	0
81				
2469 ...	4	81	169	0
81				
2470 ...	0	25	169	0
81				
2471 ...	36	196	169	0
81				

	SQBovercrowding	SQBdependency	SQBmeaned	agesq	Target
2467	18.777779	5.0625	16.0	256	1
2468	18.777779	5.0625	16.0	0	1
2469	18.777779	5.0625	16.0	81	1
2470	18.777779	5.0625	16.0	25	1
2471	18.777779	5.0625	16.0	196	1

```

[5 rows x 142 columns]

```

We have lots of string values representing house id.

```

IQ_train_df['dependency'].value_counts()

```

```

yes      2192
no       1747

```

.5	1497
2	730
1.5	713
.33333334	598
.66666669	487
8	378
.25	260
3	236
4	100
.75	98
.2	90
.40000001	84
1.3333334	84
2.5	77
5	24
.80000001	18
1.25	18
3.5	18
2.25	13
.71428573	12
1.2	11
.22222222	11
.83333331	11
1.75	11
.2857143	9
.60000002	8
1.6666666	8
.16666667	7
6	7

Name: dependency, dtype: int64

There are few numeric values with two strings as yes and no. Here we can encode yes and no as 1 and 0 respectively.

```
IQ_train_df['edjefe'].value_counts()
```

no	3762
6	1845
11	751
9	486
3	307
15	285
8	257
7	234
5	222
14	208
17	202
2	194
4	137
16	134
yes	123

```

12      113
10      111
13     103
21      43
18      19
19      14
20       7
Name: edjefe, dtype: int64

IQ_train_df['edjefa'].value_counts()

```

```

no      6230
6       947
11      399
9       237
8       217
15      188
7       179
5       176
3       152
4       136
14      120
16      113
10       96
2        84
17       76
12       72
yes       69
13       52
21        5
19        4
18        3
20        2
Name: edjefa, dtype: int64

```

Similar to dependancy, we can encode for yes and no in the above two columns.

```

encode = {'yes': 1, 'no': 0}

```

```

IQ_train_df['dependency'].replace(encode, inplace=True)
IQ_train_df['edjefe'].replace(encode, inplace=True)
IQ_train_df['edjefa'].replace(encode, inplace=True)

IQ_train_df['edjefa'].unique()

array([0, '11', '4', '10', '9', '15', '7', '14', '13', '8', '17', '6',
        '5', '3', '16', '19', 1, '21', '12', '2', '20', '18'],
      dtype=object)

```

Now lets convert all the three columns to numeric as there is no yes and no values left.

```

IQ_train_df['dependency'] = pd.to_numeric(IQ_train_df['dependency'])
IQ_train_df['edjefe'] = pd.to_numeric(IQ_train_df['edjefe'])
IQ_train_df['edjefa'] = pd.to_numeric(IQ_train_df['edjefa'])

print(IQ_train_df['dependency'].dtype)
print(IQ_train_df['edjefe'].dtype)
print(IQ_train_df['edjefa'].dtype)

float64
int64
int64

```

**Task 3: Check if there are any biases in your dataset.**

**Task 7: Count how many null values are existing in columns.**

**Task 8: Remove null value rows of the target variable.**

*# Check Target column values*

```
IQ_train_df['Target'].value_counts()
```

```
4    5996
```

```
2    1597
```

```
3    1209
```

```
1     755
```


```
Name: Target, dtype: int64
```

There are 4 poverty levels (1,2,3,4). Its an imbalanced dataset as level 4 has more counts than others.

```

sns.countplot(IQ_train_df['Target'])
plt.show()

```



Lets do NaN check.

```
col_nan_counts = IQ_train_df.isna().sum()
col_nan_counts[col_nan_counts > 0]
```

```
v2a1          6860
v18q1         7342
rez_esc       7928
meaneduc         5
SQBmeaned         5
dtype: int64
```

Above 5 columns have NaN values. Total we have 9557 records. If few columns have large number of NaN values, its impossible to imput. We can drop those columns. For the columns with very less NaN values, we can do mean imputation.

```
IQ_train_df.drop(['v2a1', 'v18q1', 'rez_esc'], axis=1, inplace=True)
```

```
IQ_train_df['meaneduc'].fillna(IQ_train_df['meaneduc'].mean(),
inplace=True)
IQ_train_df['SQBmeaned'].fillna(IQ_train_df['SQBmeaned'].mean(),
inplace=True)
```

```
IQ_train_df.columns[IQ_train_df.isna().sum() > 0]
```

```
Index([], dtype='object')
```

Now we got rid of NaN values.

Lets check for columns with constant values or zero variance.

```

IQ_train_df.drop('idhogar', axis=1).columns[IQ_train_df.var() == 0]
Index(['elimbasu5'], dtype='object')
IQ_train_df['elimbasu5'].nunique()
1

```

We can remove the 'elimbasu5' column as it has only constant value.

```

IQ_train_df.drop(['elimbasu5'], axis=1, inplace=True)
len(IQ_train_df.columns)
138

```

We have 137 columns including the target column. Its not possible to to EDA for each column and use boxplot to detect outliers.


We can select 20 columns with high corelation with target variable and do EDA for them only.

```

corr = IQ_train_df.corr()
corr_cols = corr['Target'].sort_values(ascending=False)[1:21].index
corr_cols
Index(['meaneduc', 'cielorazo', 'escolari', 'SQBescolari', 'eviv3',
      'epared3',
      'pisomoscer', 'SQBmeaned', 'paredblolad', 'etecho3',
      'SQBedjefe',
      'edjefe', 'vl8q', 'rooms', 'instlevel8', 'qmobilephone',
      'computer',
      'lugar1', 'bedrooms', 'hogar_adul'],
      dtype='object')
cols_corr = IQ_train_df[corr_cols].corr()
sns.heatmap(cols_corr)
plt.show()

```





We can see from the plot, multicollinearity is present in between few columns. So we can remove those.

```
# Delete cols with high correlation with selected columns.
def getColumnsToDelete(corr):
    cols_to_remove = set()
    for col in corr:
        if col not in cols_to_remove:
            delete_cols = corr[col][(corr[col] > 0.7) & (corr[col] <
1)].index.tolist()
            cols_to_remove.update(delete_cols)
    return cols_to_remove
```

```
cols_to_remove = getColumnsToDelete(cols_corr)
```

```
keep_cols = corr_cols[~corr_cols.isin(cols_to_remove)].tolist()
len(keep_cols)
```

16

We can do EDA with these 16 columns.

```
col_nuq = IQ_train_df[keep_cols].nunique()
col_nuq
```

```
meaneduc      156
cielorazo      2
escolari      22
eviv3          2
epared3        2
pisomoscercer  2
paredblolad    2
etecho3        2
SQBedjefe      22
vl8q           2
rooms         11
instlevel8     2
qmobilephone   11
computer       2
lugar1         2
hogar_adul     10
dtype: int64
```


We have mostly integer values representing categorical values and float values representing continuous value.

```
def getHist(col):
    sns.histplot(IQ_train_df[col], kde=True)
    plt.show()
```


```
def getBox(col):
    sns.boxplot(IQ_train_df[col])
    plt.show()
```

Column: meaneduc EDA

```
getHist('meaneduc')
```




```
getBox('meaneduc')
```



There are lots of outliers. We can set limit as 19.


```
IQ_train_df = IQ_train_df[IQ_train_df['meaneduc'] <= 19]
```

```
getBox('meaneduc')
```




Column: escolar\_i EDA

`getHist('escolar_i')`




`getBox('escolar_i')`




There is no outlier.

Column: SQBedjefe EDA

`getHist('SQBedjefe')`




`getBox('SQBedjefe')`




There are few outliers. We can set limit as 200.

```
IQ_train_df = IQ_train_df[IQ_train_df['SQBedjefe'] <= 200]
```

```
getBox('SQBedjefe')
```




```
IQ_train_df = IQ_train_df[IQ_train_df['SQBedjefe'] <= 90]  
getBox('SQBedjefe')
```




Outliers got removed by setting 90 as limit.

Column: rooms EDA

```
getHist('rooms')
```



```
getBox('rooms')
```




There are two outliers. So we can set 9 as limit.

```
IQ_train_df = IQ_train_df[IQ_train_df['rooms'] <= 9]
```


```
getBox('rooms')
```






Column: qmobilephone EDA

`getHist('qmobilephone')`



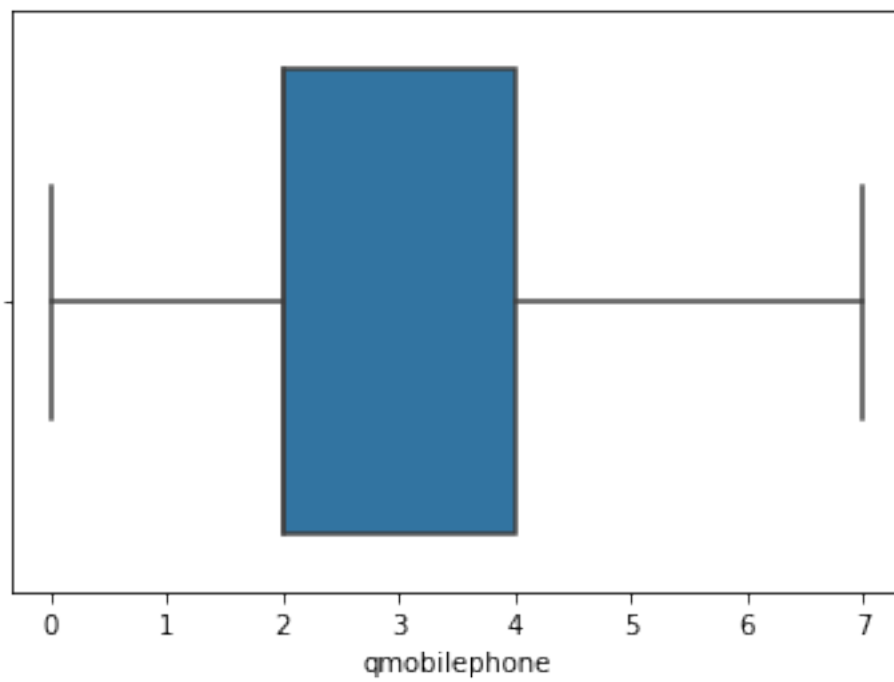
`getBox('qmobilephone')`



We can set 7 as limit here to remove the outliers.

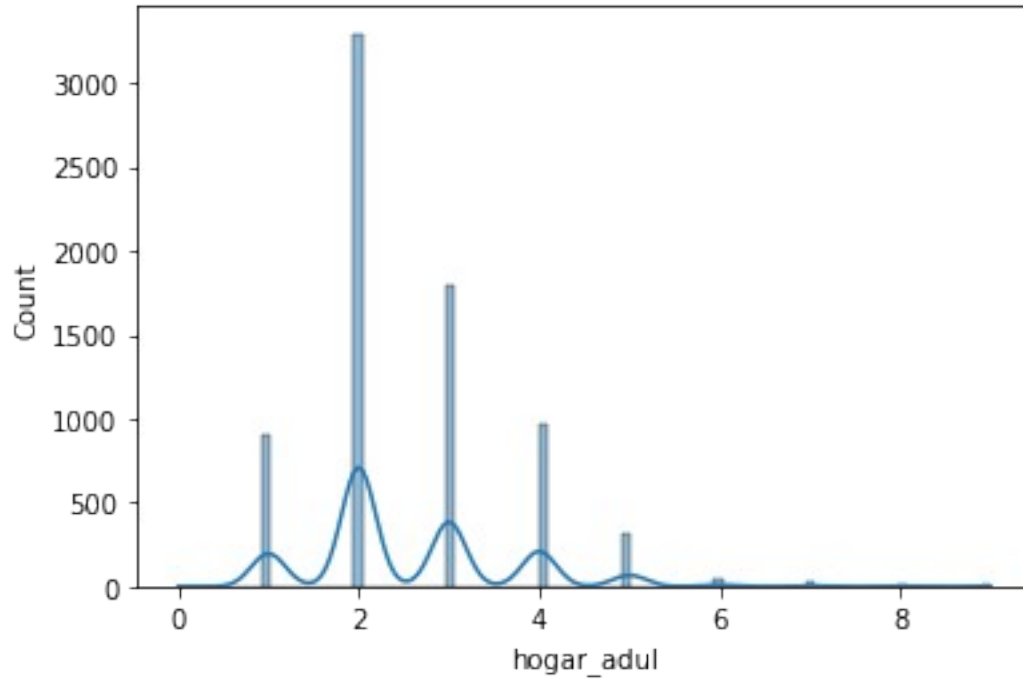
```
IQ_train_df = IQ_train_df[IQ_train_df['qmobilephone'] <= 7]
```

```
getBox('qmobilephone')
```




Column: hogar\_adul EDA

```
getHist('hogar_adul')
```



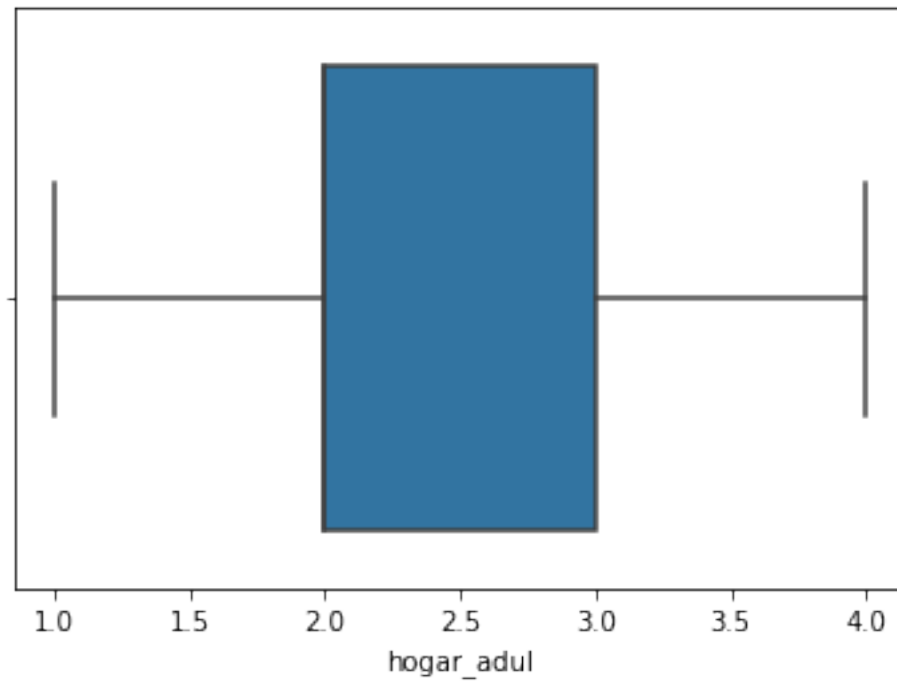
```
getBox('hogar_adul')
```



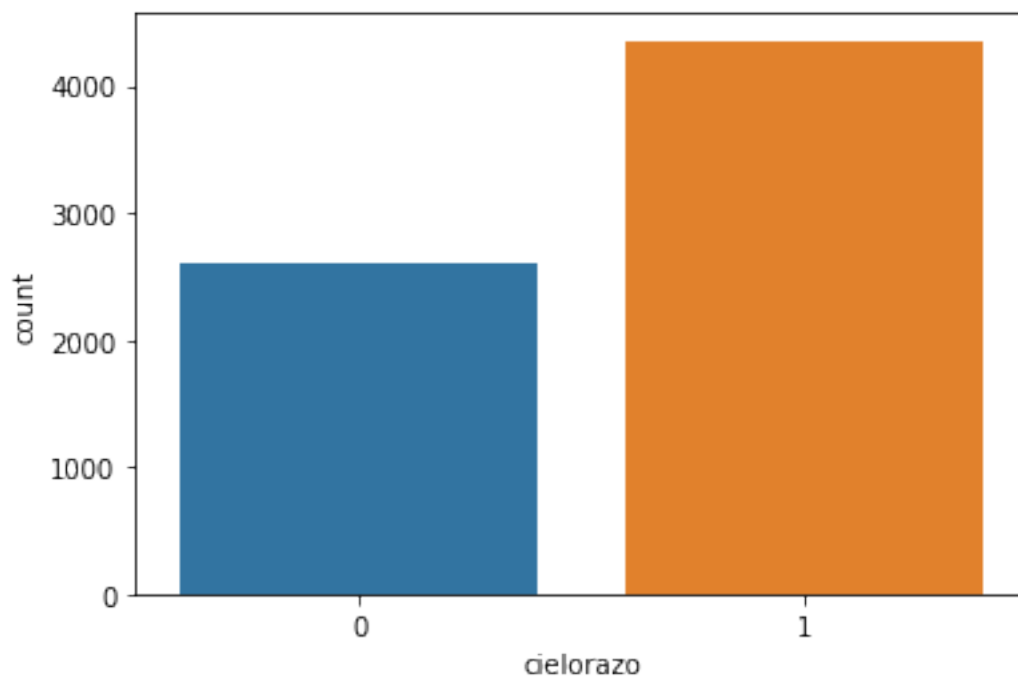
Here we need to set lower and upper limit as 1 and 4 respectively to get rid of outliers.



```
IQ_train_df = IQ_train_df[(IQ_train_df['hogar_adul'] >= 1) &  
(IQ_train_df['hogar_adul'] <= 4)]
```

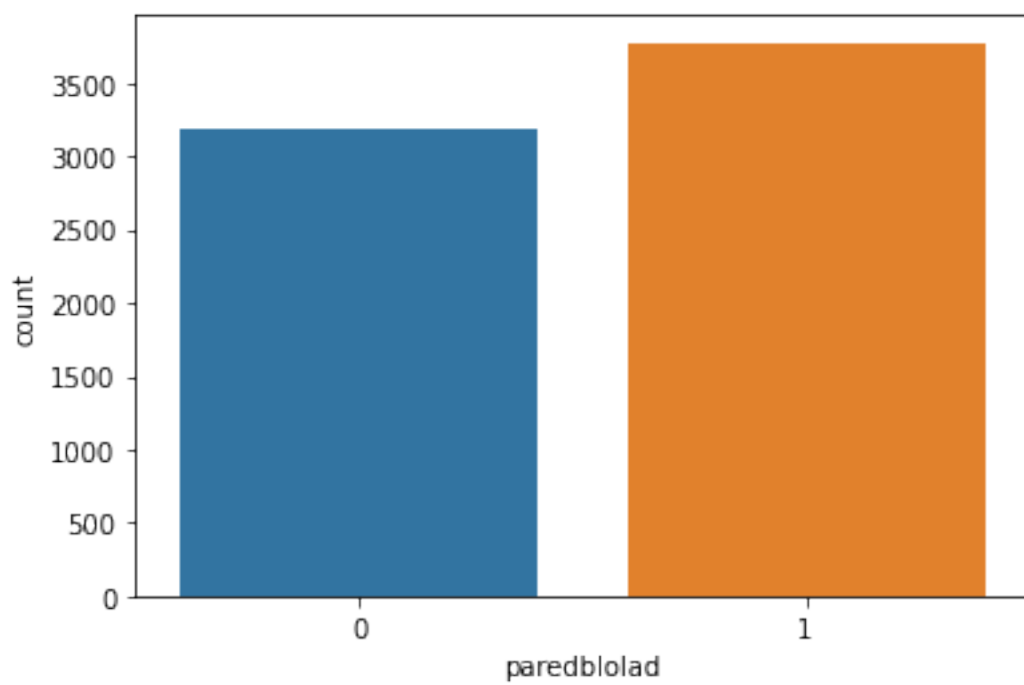

```
getBox('hogar_adul')
```






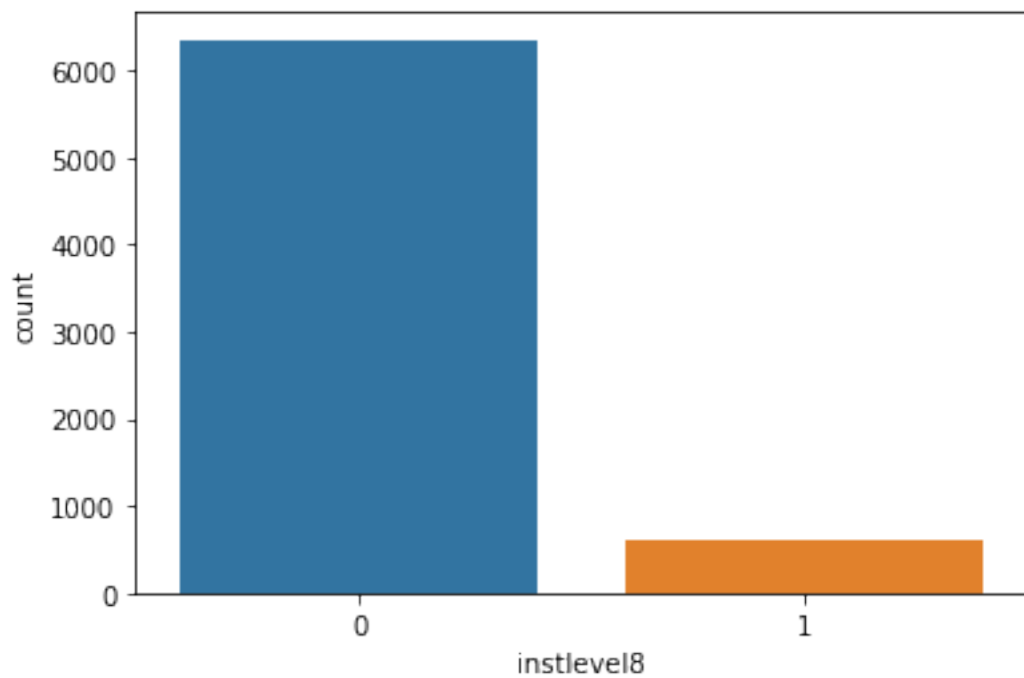
```
for index, val in zip(col_nuq.index, col_nuq):  
    if val == 2:  
        sns.countplot(IQ_train_df[index])  
        plt.show()
```
















We have reasonable values in the above columns, where for few columns, there are imbalance in its type. But its fine, we can keep it.

```
print('Final train data shape', IQ_train_df.shape)
```

Final train data shape (6963, 138)

Still we have enough records to train models.

**Task 4: Check whether all members of the house have the same poverty level.**

```
per_house_uq_level = IQ_train_df.groupby('idhogar')
['Target'].nunique()
```

```
per_house_uq_level[per_house_uq_level == 1].size /
per_house_uq_level.size
```

0.9644602398933807

Only 96.44% household have all its members with same poverty level.

```
per_house_uq_level[per_house_uq_level > 1].size
```

80

There are 80 household like below which have different poverty level for each member.

```
per_house_uq_level[per_house_uq_level > 1][:1]
```

```
idhogar
0172ab1d9    2
Name: Target, dtype: int64
```

```
IQ_train_df[IQ_train_df['idhogar'] == '0172ab1d9']
```

	hacdor	rooms	hacapo	v14a	refrig	v18q	r4h1	r4h2	r4h3
r4m1 ... \									
7651	0	5	0	1	1	0	0	2	2
0 ...									
7652	0	5	0	1	1	0	0	2	2
0 ...									
7653	0	5	0	1	1	0	0	2	2
0 ...									
7654	0	5	0	1	1	0	0	2	2
0 ...									
7655	0	5	0	1	1	0	0	2	2
0 ...									

	SQBescolari	SQBage	SQBhogar_total	SQBedjefe	SQBhogar_nin	\
7651	49	196	25	36		4
7652	100	289	25	36		4
7653	36	2601	25	36		4
7654	36	2304	25	36		4
7655	121	441	25	36		4

	SQBovercrowding	SQBdependency	SQBmeand	agesq	Target
7651	2.777778	0.444444	58.777775	196	3
7652	2.777778	0.444444	58.777775	289	2
7653	2.777778	0.444444	58.777775	2601	3
7654	2.777778	0.444444	58.777775	2304	3
7655	2.777778	0.444444	58.777775	441	2

[5 rows x 138 columns]

For the household '0172ab1d9', members dont have same poverty level.

#### Task 5: Check if there is a house without a family head.

```
household_head = IQ_train_df.groupby('idhogar')['parentesco1'].sum()
household_head[household_head == 0].index
```

```
Index(['03c6bdf85', '09b195e7a', '1bc617b23', '374ca5a19',
      '61c10e099',
      '6b1b2405f', '896fe6d3e', 'a0812ef17', 'ad687ad89',
      'b1f4d89d7',
      'bfd5067c2', 'd363d9183', 'f2bfa75c4'],
      dtype='object', name='idhogar')
```

Above are the houses with no family head.

#### Task 6: Set poverty level of the members and the head of the house within a family.

```
houseId_mismatch_level = per_house_uq_level[per_house_uq_level >
1].index.tolist()
len(houseId_mismatch_level)
```

80

There are 80 houses with different poverty level of each member. We need to update the level of each member same as the head i a house.

```
houseId_mismatch_level[0]
```

```
'0172ab1d9'
```

```
IQ_train_df[IQ_train_df['idhogar'] == '0172ab1d9']
```

	hacdor	rooms	hacapo	v14a	refrig	v18q	r4h1	r4h2	r4h3
r4m1 ... \									
7651	0	5	0	1	1	0	0	2	2
0 ...									
7652	0	5	0	1	1	0	0	2	2
0 ...									
7653	0	5	0	1	1	0	0	2	2
0 ...									
7654	0	5	0	1	1	0	0	2	2
0 ...									
7655	0	5	0	1	1	0	0	2	2
0 ...									

	SQBescolari	SQBage	SQBhogar_total	SQBedjefe	SQBhogar_nin	\
7651	49	196	25	36		4
7652	100	289	25	36		4
7653	36	2601	25	36		4
7654	36	2304	25	36		4
7655	121	441	25	36		4

	SQBovercrowding	SQBdependency	SQBmeaned	agesq	Target
7651	2.777778	0.444444	58.777775	196	3
7652	2.777778	0.444444	58.777775	289	2
7653	2.777778	0.444444	58.777775	2601	3
7654	2.777778	0.444444	58.777775	2304	3
7655	2.777778	0.444444	58.777775	441	2

```
[5 rows x 138 columns]
```

```
for houseId in houseId_mismatch_level:
    head_level = IQ_train_df[(IQ_train_df['idhogar'] == '0172ab1d9') &
(IQ_train_df['parentesco1'] == 1)]['Target'].values[0]
    IQ_train_df.loc[IQ_train_df['idhogar'] == houseId, 'Target'] =
head_level
```

```
per_house_uq_level = IQ_train_df.groupby('idhogar')
['Target'].nunique()
per_house_uq_level[per_house_uq_level > 1].size
```

```
0
```

Now all the household have same poverty level.

**Task 9: Predict the accuracy using random forest classifier.**

```
x = IQ_train_df.drop(['Target'], axis=1)
y = IQ_train_df['Target']
```

```
print(x.shape)
print(y.shape)
```

```
(6963, 137)
(6963,)
```

Now we dont need the houseId (idhogar), Lets delete that column too.

```
x.drop(['idhogar'], axis=1, inplace=True)
```

```
print(x.shape)
```

```
(6963, 136)
```

Before proceeding with the modeling, lets get rid of columns with high multi-collinearity.

```
colsToDelete = getColumnsToDelete(x.corr())
len(colsToDelete)
```

```
20
```

We can remove 20 columns.

```
x = x[x.columns[~x.columns.isin(colsToDelete)]]
```

```
print(x.shape)
```

```
(6963, 116)
```

*# Do train test split*

```
x_train, x_test, y_train, y_test = train_test_split(x, y)
```

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(5222, 116)
(1741, 116)
(5222,)
(1741,)
```

*# Create a RandomForestClassifier object*

```
rfc = RandomForestClassifier(class_weight='balanced') # As we have  
imbalanced dataset, we use class_weight
```

```
# Get accuracy
```

```
rfc.fit(x_train, y_train)  
y_pred = rfc.predict(x_test)
```

```
print('Accuracy score for RandomForestClassifier: ',  
accuracy_score(y_test, y_pred))
```


Accuracy score for RandomForestClassifier: 0.9310740953475014

```
cols_idx = rfc.feature_importances_.argsort()[::-1][:10]
```

```
important_features = x.columns[cols_idx]
```

```
importance_values = rfc.feature_importances_[cols_idx]
```

```
plt.figure(figsize=(10, 8))  
sns.barplot(important_features, importance_values)  
plt.xticks(rotation = 90)  
plt.show()
```



### Task 10: Check the accuracy using random forest with cross validation.

Lets try out cross validation with GridSearchCV to find best parameter for RandomForestClassifier

```
params = {'n_estimators': [100, 150, 200], 'max_depth': [5, 10, 15],  
'min_samples_split': [2, 4, 6]}
```

```
rfc_gscv = RandomForestClassifier(class_weight='balanced')  
gscv = GridSearchCV(rfc_gscv, params, cv = 5, scoring='accuracy')
```

```
gscv.fit(x_train, y_train)
```

```
GridSearchCV(cv=5,  
estimator=RandomForestClassifier(class_weight='balanced'),  
param_grid={'max_depth': [5, 10, 15],  
            'min_samples_split': [2, 4, 6],  
            'n_estimators': [100, 150, 200]},  
scoring='accuracy')
```

```
best_rfc = gscv.best_estimator_
```

```
best_rfc
```

```
RandomForestClassifier(class_weight='balanced', max_depth=15,  
n_estimators=200)
```

```
gscv.best_params_
```

```
{'max_depth': 15, 'min_samples_split': 2, 'n_estimators': 200}
```

```
print('Accuracy with best classifier', gscv.best_score_)
```

```
Accuracy with best classifier 0.8923809785697262
```

Lets train the model and test it.

```
best_rfc.fit(x_train, y_train)  
y_pred = best_rfc.predict(x_test)
```

```
print('Accuracy score for RandomForestClassifier: ',  
accuracy_score(y_test, y_pred))
```

```
Accuracy score for RandomForestClassifier: 0.9247558874210224
```

Accuracy is almost the same.

Lets use the test dataset to predict.

```
IQ_test_df = pd.read_csv('test.csv')
```

```
IQ_test_df.head()
```

	Id	v2a1	hacdor	rooms	hacapo	v14a	refrig	v18q
v18q1 \								
0	ID_2f6873615	NaN	0	5	0	1	1	0
NaN								
1	ID_1c78846d2	NaN	0	5	0	1	1	0
NaN								
2	ID_e5442cf6a	NaN	0	5	0	1	1	0
NaN								
3	ID_a8db26a79	NaN	0	14	0	1	1	1
1.0								
4	ID_a62966799	175000.0	0	4	0	1	1	1
1.0								

	r4h1	...	age	SQBescolari	SQBage	SQBhogar_total	SQBedjefe	\
0	1	...	4	0	16	9	0	
1	1	...	41	256	1681	9	0	
2	1	...	41	289	1681	9	0	
3	0	...	59	256	3481	1	256	
4	0	...	18	121	324	1	0	

	SQBhogar_nin	SQBovercrowding	SQBdependency	SQBmeand	agesq
0	1	2.25	0.25	272.25	16
1	1	2.25	0.25	272.25	1681
2	1	2.25	0.25	272.25	1681
3	0	1.00	0.00	256.00	3481
4	1	0.25	64.00	NaN	324

[5 rows x 142 columns]

```
# Drop columns
```

```
colsToDelete.update(['Id', 'idhogar', 'elimbasu5', 'v2a1', 'v18q1', 'rez_esc'])
```

```
IQ_test_df.drop(columns=colsToDelete, inplace=True, axis=1)
```

```
print(IQ_test_df.shape)
```

```
(23856, 116)
```

```
IQ_test_df['dependency'].replace(encode, inplace=True)
```

```
IQ_test_df['edjefe'].replace(encode, inplace=True)
```

```
IQ_test_df['edjefa'].replace(encode, inplace=True)
```

```
IQ_test_df['dependency'] = pd.to_numeric(IQ_test_df['dependency'])
```

```
IQ_test_df['edjefe'] = pd.to_numeric(IQ_test_df['edjefe'])
```

```
IQ_test_df['edjefa'] = pd.to_numeric(IQ_test_df['edjefa'])
```

```
IQ_test_df.dropna(inplace=True)
```

```
best_rfc.predict(IQ_test_df)
```

```
array([4, 4, 4, ..., 3, 3, 3])
```