

## Lending Club Loan Data Analysis

```
# import libraries
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, BatchNormalization, Dropout
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.metrics import confusion_matrix, auc, roc_curve
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
# Get dataset
```

```
loan_df = pd.read_csv('loan_data.csv')
loan_df.head()
```

	credit.policy		purpose	int.rate	installment
log.annual.inc \					
0	1	debt_consolidation	0.1189	829.10	
11.350407					
1	1	credit_card	0.1071	228.22	
11.082143					
2	1	debt_consolidation	0.1357	366.86	
10.373491					
3	1	debt_consolidation	0.1008	162.34	
11.350407					
4	1	credit_card	0.1426	102.92	
11.299732					

	dti	fico	days.with.cr.line	revol.bal	revol.util
inq.last.6mths \					
0	19.48	737	5639.958333	28854	52.1
0					
1	14.29	707	2760.000000	33623	76.7
0					
2	11.63	682	4710.000000	3511	25.6
1					
3	8.10	712	2699.958333	33667	73.2
1					
4	14.97	667	4066.000000	4740	39.5
0					

```
delinq.2yrs  pub.rec  not.fully.paid
```

0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	1	0	0

```
loan_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   credit.policy          9578 non-null   int64
1   purpose                9578 non-null   object
2   int.rate               9578 non-null   float64
3   installment            9578 non-null   float64
4   log.annual.inc         9578 non-null   float64
5   dti                    9578 non-null   float64
6   fico                   9578 non-null   int64
7   days.with.cr.line      9578 non-null   float64
8   revol.bal              9578 non-null   int64
9   revol.util             9578 non-null   float64
10  inq.last.6mths         9578 non-null   int64
11  delinq.2yrs            9578 non-null   int64
12  pub.rec                9578 non-null   int64
13  not.fully.paid         9578 non-null   int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

As we can see here, there is no NaN value in any column.

```
loan_df.shape
```

```
(9578, 14)
```

From data, we can say, we have few categorical columns like credit.policy, purpose and columns with limited range of values like inq.last.6mths, delinq.2yrs, pub.rec and not.fully.paid.

```
loan_df.describe()
```

	credit.policy	int.rate	installment	log.annual.inc
dti \				
count	9578.000000	9578.000000	9578.000000	9578.000000
mean	0.804970	0.122640	319.089413	10.932117
std	0.396245	0.026847	207.071301	0.614813
min	0.000000	0.060000	15.670000	7.547502
max	12.606679	6.883970	9578.000000	10.932117

25%	1.000000	0.103900	163.770000	10.558414
7.212500				
50%	1.000000	0.122100	268.950000	10.928884
12.665000				
75%	1.000000	0.140700	432.762500	11.291293
17.950000				
max	1.000000	0.216400	940.140000	14.528354
29.960000				

	fico	days.with.cr.line	revol.bal	revol.util \
count	9578.000000	9578.000000	9.578000e+03	9578.000000
mean	710.846314	4560.767197	1.691396e+04	46.799236
std	37.970537	2496.930377	3.375619e+04	29.014417
min	612.000000	178.958333	0.000000e+00	0.000000
25%	682.000000	2820.000000	3.187000e+03	22.600000
50%	707.000000	4139.958333	8.596000e+03	46.300000
75%	737.000000	5730.000000	1.824950e+04	70.900000
max	827.000000	17639.958330	1.207359e+06	119.000000



	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid
count	9578.000000	9578.000000	9578.000000	9578.000000
mean	1.577469	0.163708	0.062122	0.160054
std	2.200245	0.546215	0.262126	0.366676
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000	0.000000
75%	2.000000	0.000000	0.000000	0.000000
max	33.000000	13.000000	5.000000	1.000000

#### EDA Of numeric data



```
def hist_box_plot(col):
    plt.figure(figsize=(12, 5))
    plt.subplot(1,2,1)
    sns.histplot(data=loan_df, x= col, hue='credit.policy', kde=True)
    plt.subplot(1,2,2)
    sns.boxplot(loan_df[col])
    plt.show()

def box_plot(df, col):
    sns.boxplot(df[col])
    plt.show()



hist_box_plot('int.rate')
```





```
hist_box_plot('installment')
```





```
hist_box_plot('log.annual.inc')
```





```
hist_box_plot('dti')
```





`hist_box_plot('fico')`




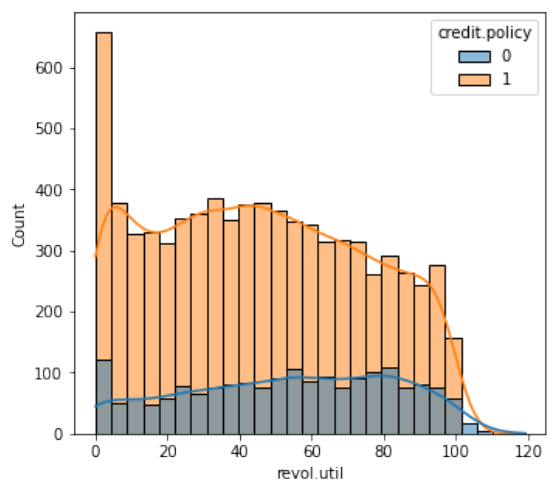
`hist_box_plot('days.with.cr.line')`



`hist_box_plot('revol.bal')`



`hist_box_plot('revol.util')`




We are getting sensible distribution plots for the numeric values. Also for customer meeting credit criteria, all the numeric column values are high.

We can remove the outlier by considering minimum data loss (max 5% data loss).


```
loan_df[loan_df['int.rate'] < 0.19].shape[0]/ loan_df.shape[0] * 100
99.01858425558572
```

```
filter_loan_df = loan_df[loan_df['int.rate'] < 0.19]
```

```
box_plot(filter_loan_df, 'int.rate')
```



```
loan_df = filter_loan_df  
loan_df[loan_df['installment'] < 750].shape[0]/ loan_df.shape[0] * 100  
95.08646140868832  
filter_loan_df = loan_df[loan_df['installment'] < 750]  
box_plot(filter_loan_df, 'installment')
```



```

loan_df = filter_loan_df


loan_df[(loan_df['log.annual.inc'] > 9.5) & (loan_df['log.annual.inc']
< 12.3)].shape[0]/ loan_df.shape[0] * 100

97.26103348857839

filter_loan_df = loan_df[(loan_df['log.annual.inc'] > 9.5) &
(loan_df['log.annual.inc'] < 12.3)]

box_plot(filter_loan_df, 'log.annual.inc')

```



```

loan_df = filter_loan_df

loan_df[loan_df['fico'] < 820].shape[0]/ loan_df.shape[0] * 100


99.94299395735948

filter_loan_df = loan_df[loan_df['fico'] < 820]


box_plot(filter_loan_df, 'fico')

```





```
loan_df = filter_loan_df  
loan_df[loan_df['days.with.cr.line'] < 9300].shape[0]/  
loan_df.shape[0] * 100  
95.25439196897102  
filter_loan_df = loan_df[loan_df['days.with.cr.line'] < 9300]  
box_plot(filter_loan_df, 'days.with.cr.line')
```



```
loan_df = filter_loan_df
```

```
loan_df[loan_df['revol.bal'] < 0.05].shape[0] / loan_df.shape[0] * 100
```

```
3.1137724550898205
```

We cant remove the outlier here, number of outliers are so high.

Now lets check the shape.


```
loan_df.shape
```

```
(8350, 14)
```

*EDA on categorical data & numeric with limited number of values*


```
def count_plot(col, hue_val= None, rotate= False):
    sns.countplot(data= loan_df,
                  x= col,
                  hue= hue_val,
                  order = loan_df[col].value_counts().index)
    if rotate:
        plt.xticks(rotation=90)
```

```
count_plot('credit.policy')
```



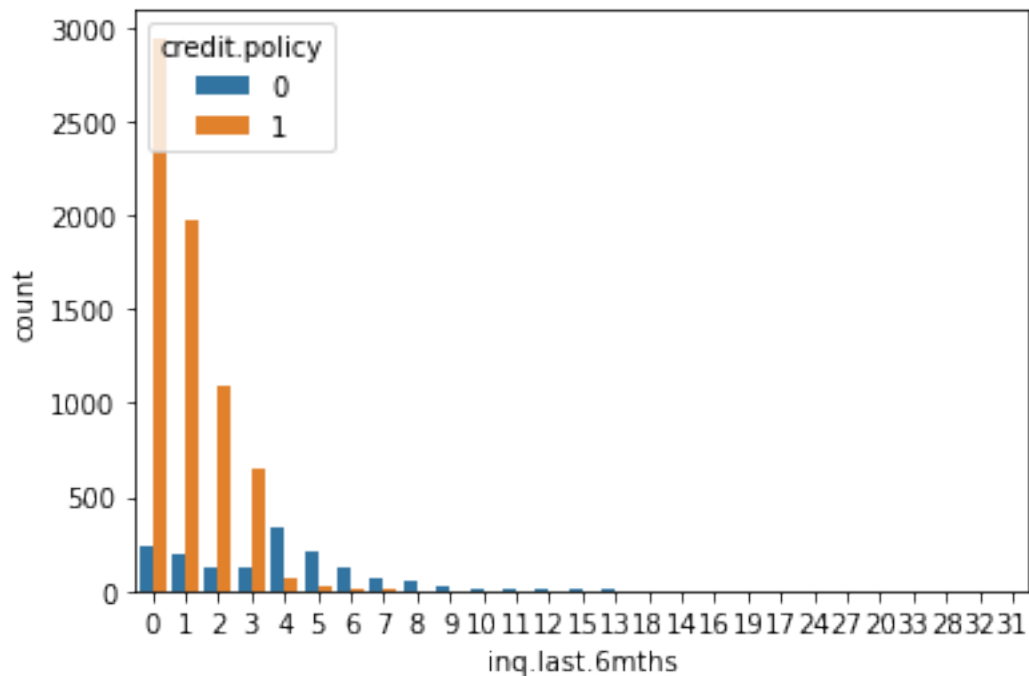
Its a imbalanced data.

```
count_plot('purpose', hue_val='credit.policy', rotate= True)
```



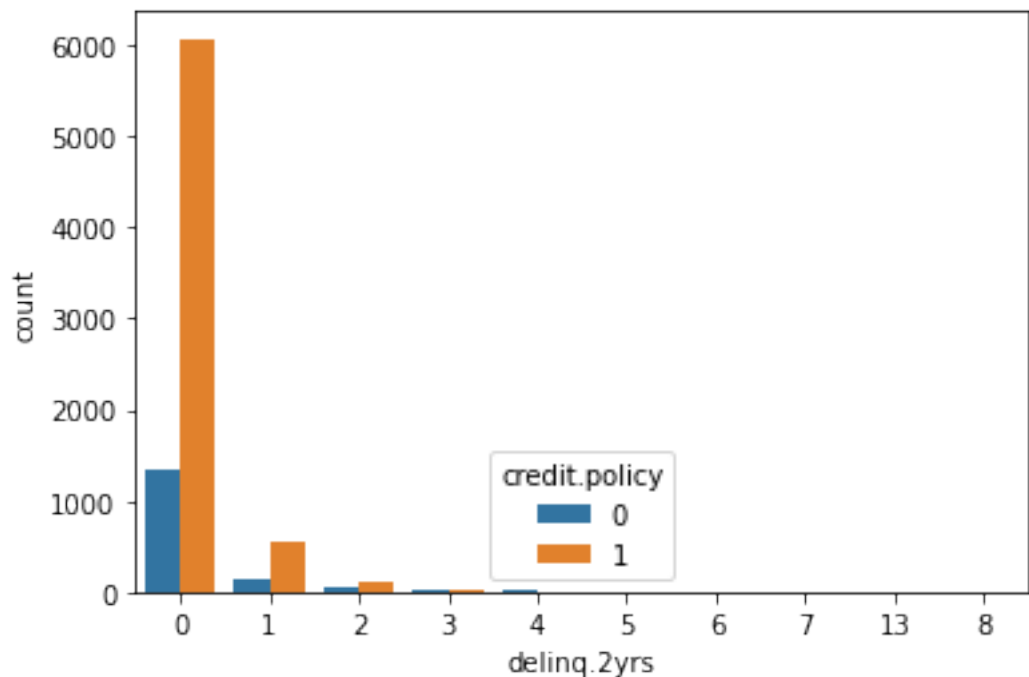
Most of the loans are applied for debt\_consolidation where as educational loan count is minimum.

```
count_plot('inq.last.6mths', hue_val='credit.policy')
```



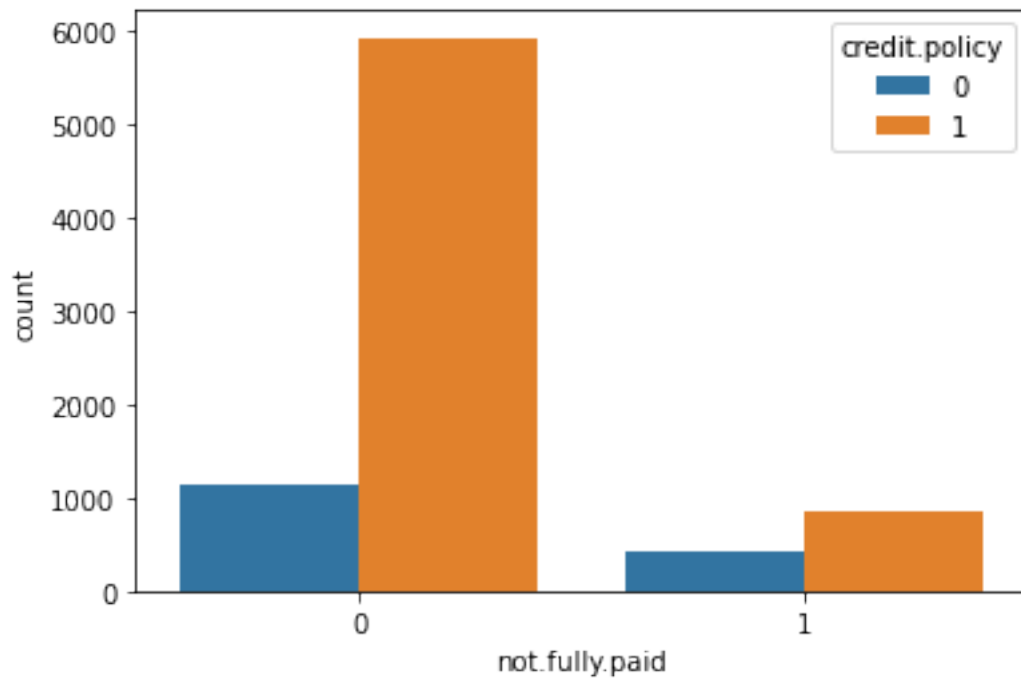
Most of the customers have 0-3 number of queries in last 6 months.

```
count_plot('delinq.2yrs', hue_val='credit.policy')
```



There are very few people with higher number of past dues.

```
count_plot('not.fully.paid', hue_val='credit.policy')
```



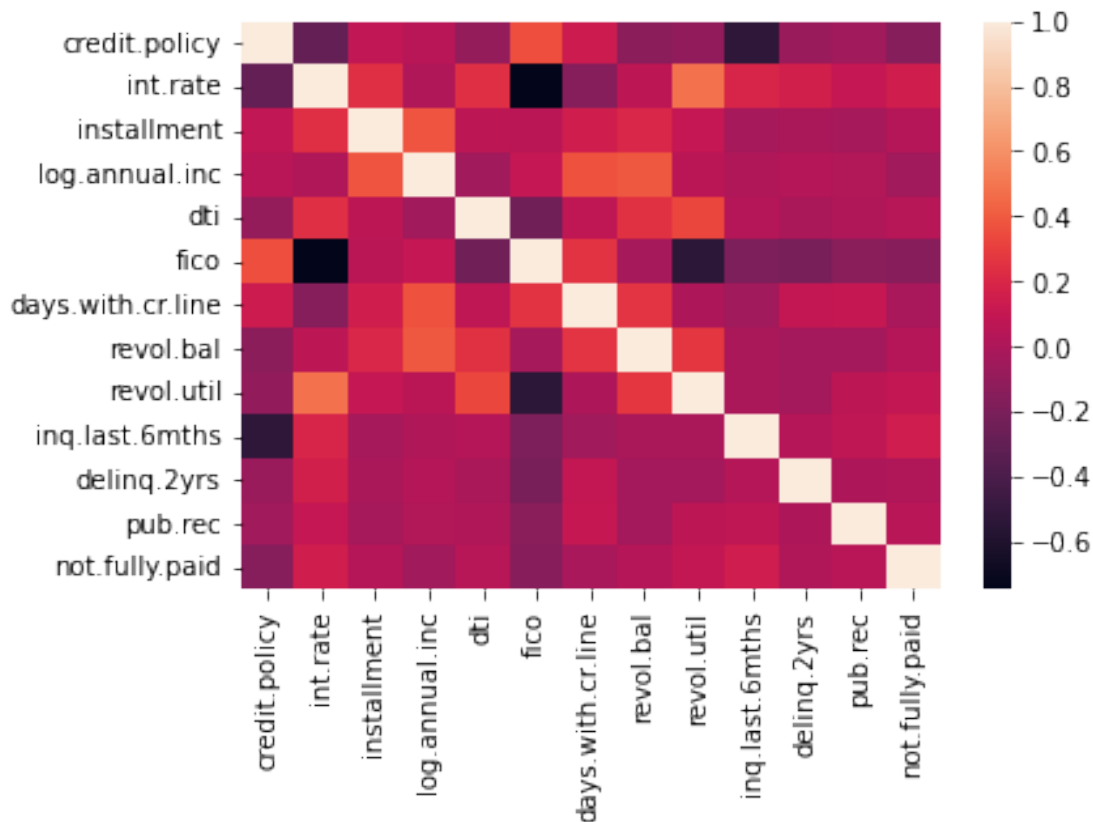
There are many customers who have fully paid and are eligible for credit.

*Additional Feature Engineering (Remove columns with high correlation if any with other feature columns)*

```
loan_corr = loan_df.corr()
```

```
sns.heatmap(loan_corr)
```

```
plt.show()
```



There are many columns which are correlated with each other. As we have 13 feature columns and by looking at the column names, most of them are important features and impact the policy credit.

But we can filter out the columns with correlation among each other.

```
target_corr_cols =
loan_corr['credit.policy'].sort_values(ascending=False)
[1:].index.tolist()
target_corr_cols

['fico',
'days.with.cr.line',
'installment',
'log.annual.inc',
'pub.rec',
'delinq.2yrs',
'dti',
'revol.util',
'revol.bal',
'not.fully.paid',
'int.rate',
'inq.last.6mths']
```

```

for col in target_corr_cols:
    corr_vals = loan_corr[col]
    corr_cols = corr_vals[corr_vals >= 0.6].index.tolist()
    print(col, corr_cols)

```

```

fico ['fico']
days.with.cr.line ['days.with.cr.line']
installment ['installment']
log.annual.inc ['log.annual.inc']
pub.rec ['pub.rec']
delinq.2yrs ['delinq.2yrs']
dti ['dti']
revol.util ['revol.util']
revol.bal ['revol.bal']
not.fully.paid ['not.fully.paid']
int.rate ['int.rate']
inq.last.6mths ['inq.last.6mths']

```

As we can see, we have set the high correlation lower limit as 0.6 and no columns are correlated with each other.

So we can use all the columns for modelling.

Now lets encode the purpose column and get train and test dataset out of the loan\_df.

*Feature Transformation (Convert categorical values to numeric vectors)*

```

loan_df_en = pd.get_dummies(data= loan_df, columns=['purpose'],
drop_first=True, dtype= 'int64')

```

```

loan_df_en.shape

```

```

(8350, 19)

```

```

loan_df_en.head()

```

	credit.policy	int.rate	installment	log.annual.inc	dti	
fico \						
1	1	0.1071	228.22	11.082143	14.29	707
2	1	0.1357	366.86	10.373491	11.63	682
3	1	0.1008	162.34	11.350407	8.10	712
4	1	0.1426	102.92	11.299732	14.97	667
5	1	0.0788	125.13	11.904968	16.98	727

	days.with.cr.line	revol.bal	revol.util	inq.last.6mths
delinq.2yrs \				
1	2760.000000	33623	76.7	0



0				
2	4710.000000	3511	25.6	1
0				
3	2699.958333	33667	73.2	1
0				
4	4066.000000	4740	39.5	0
1				
5	6120.041667	50807	51.0	0
0				

	pub.rec	not.fully.paid	purpose_credit_card
purpose_debt_consolidation \			
1	0	0	1
0			
2	0	0	0
1			
3	0	0	0
1			
4	0	0	1
0			
5	0	0	1
0			

	purpose_educational	purpose_home_improvement
purpose_major_purchase \		
1	0	0
0		
2	0	0
0		
3	0	0
0		
4	0	0
0		
5	0	0
0		

	purpose_small_business
1	0
2	0
3	0
4	0
5	0

loan\_df\_en.columns

```
Index(['credit.policy', 'int.rate', 'installment', 'log.annual.inc',
      'dti',
      'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
      'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid',
      'purpose_credit_card', 'purpose_debt_consolidation',
```

```
        'purpose_educational', 'purpose_home_improvement',  
        'purpose_major_purchase', 'purpose_small_business'],  
        dtype='object')
```

#### *Preprocessing steps for modelling*

```
x = loan_df_en.drop(['credit.policy'], axis=1)  
y = loan_df_en['credit.policy']
```

```
print(x.shape)  
print(y.shape)
```

```
(8350, 18)  
(8350,)
```

#### *# train test split*

```
x_train, x_test, y_train, y_test = train_test_split(x, y)
```

```
print(x_train.shape)  
print(x_test.shape)  
print(y_train.shape)  
print(y_test.shape)
```

```
(6262, 18)  
(2088, 18)  
(6262,)  
(2088,)
```

#### *# Standardize the values*

```
ss = StandardScaler()
```

```
x_train = ss.fit_transform(x_train)  
x_test = ss.transform(x_test)
```

#### *Deep Learning Model*

```
model = Sequential()  
model.add(Dense(128, activation= 'relu', input_shape=(18,),  
                kernel_initializer='he_normal',  
                kernel_regularizer='L2'))  
model.add(Dropout(0.5))  
model.add(BatchNormalization())  
model.add(Dense(64, activation= 'relu',  
                kernel_initializer='he_normal',  
                kernel_regularizer='L2'))  
model.add(Dropout(0.5))  
model.add(BatchNormalization())  
model.add(Dense(32, activation= 'relu',  
                kernel_initializer='he_normal',  
                kernel_regularizer='L2'))  
model.add(BatchNormalization())
```

```

model.add(Dense(10, activation= 'relu',
                kernel_initializer='he_normal',
                kernel_regularizer='L2'))
model.add(Dense(1, activation= 'sigmoid',
                kernel_initializer='glorot_normal',
                kernel_regularizer='L2'))

```

```
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	2432
dropout (Dropout)	(None, 128)	0
batch_normalization (Batch Normalization)	(None, 128)	512
dense_1 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 64)	256
dense_2 (Dense)	(None, 32)	2080
batch_normalization_2 (Batch Normalization)	(None, 32)	128
dense_3 (Dense)	(None, 10)	330
dense_4 (Dense)	(None, 1)	11

```

Total params: 14,005
Trainable params: 13,557
Non-trainable params: 448

```

```

model.compile(optimizer='adam',
              loss=tf.keras.losses.binary_crossentropy,
              metrics=[tf.keras.metrics.AUC()])

```

```
callback = EarlyStopping(monitor='val_loss', patience=5)
```

```

history = model.fit(
x_train, y_train,
validation_data=(x_test, y_test),

```

```
batch_size= 30,  
epochs= 20,  
verbose=1,  
callbacks=[callback]  
)
```

Epoch 1/20

209/209 [=====] - 3s 5ms/step - loss: 4.2353  
- auc: 0.6486 - val\_loss: 3.1175 - val\_auc: 0.8656

Epoch 2/20

209/209 [=====] - 1s 4ms/step - loss: 2.4746  
- auc: 0.8194 - val\_loss: 1.8458 - val\_auc: 0.9184

Epoch 3/20

209/209 [=====] - 1s 5ms/step - loss: 1.5379  
- auc: 0.8478 - val\_loss: 1.1667 - val\_auc: 0.9327

Epoch 4/20

209/209 [=====] - 1s 5ms/step - loss: 1.0066  
- auc: 0.8807 - val\_loss: 0.7855 - val\_auc: 0.9406

Epoch 5/20

209/209 [=====] - 1s 5ms/step - loss: 0.7113  
- auc: 0.8983 - val\_loss: 0.5892 - val\_auc: 0.9428

Epoch 6/20

209/209 [=====] - 1s 4ms/step - loss: 0.5565  
- auc: 0.9007 - val\_loss: 0.4831 - val\_auc: 0.9459

Epoch 7/20

209/209 [=====] - 1s 4ms/step - loss: 0.4541  
- auc: 0.9160 - val\_loss: 0.4052 - val\_auc: 0.9462

Epoch 8/20

209/209 [=====] - 1s 5ms/step - loss: 0.4077  
- auc: 0.9166 - val\_loss: 0.3703 - val\_auc: 0.9543

Epoch 9/20

209/209 [=====] - 1s 4ms/step - loss: 0.3748  
- auc: 0.9210 - val\_loss: 0.3576 - val\_auc: 0.9552

Epoch 10/20

209/209 [=====] - 1s 5ms/step - loss: 0.3508  
- auc: 0.9292 - val\_loss: 0.3290 - val\_auc: 0.9511

Epoch 11/20

209/209 [=====] - 1s 4ms/step - loss: 0.3407  
- auc: 0.9311 - val\_loss: 0.3127 - val\_auc: 0.9508

Epoch 12/20

209/209 [=====] - 1s 4ms/step - loss: 0.3374  
- auc: 0.9289 - val\_loss: 0.3027 - val\_auc: 0.9581

Epoch 13/20

209/209 [=====] - 1s 4ms/step - loss: 0.3189  
- auc: 0.9384 - val\_loss: 0.2777 - val\_auc: 0.9645

Epoch 14/20

209/209 [=====] - 1s 4ms/step - loss: 0.3238  
- auc: 0.9337 - val\_loss: 0.2697 - val\_auc: 0.9662

Epoch 15/20

209/209 [=====] - 1s 4ms/step - loss: 0.3104

```

- auc: 0.9414 - val_loss: 0.2646 - val_auc: 0.9644
Epoch 16/20
209/209 [=====] - 1s 4ms/step - loss: 0.3109
- auc: 0.9375 - val_loss: 0.2971 - val_auc: 0.9725
Epoch 17/20
209/209 [=====] - 1s 3ms/step - loss: 0.2951
- auc: 0.9470 - val_loss: 0.2688 - val_auc: 0.9708
Epoch 18/20
209/209 [=====] - 1s 3ms/step - loss: 0.2952
- auc: 0.9449 - val_loss: 0.2514 - val_auc: 0.9747
Epoch 19/20
209/209 [=====] - 1s 3ms/step - loss: 0.2851
- auc: 0.9504 - val_loss: 0.2423 - val_auc: 0.9715
Epoch 20/20
209/209 [=====] - 1s 3ms/step - loss: 0.2942
- auc: 0.9459 - val_loss: 0.2454 - val_auc: 0.9756

```

```

history = pd.DataFrame(history.history)
history.head()


```

	loss	auc	val_loss	val_auc
0	4.235277	0.648598	3.117543	0.865618
1	2.474607	0.819444	1.845755	0.918365
2	1.537950	0.847843	1.166706	0.932721
3	1.006636	0.880666	0.785537	0.940568
4	0.711283	0.898319	0.589222	0.942774

```

plt.figure(figsize = (15,6))
plt.plot(history.val_loss, label='Test Set Loss')
plt.plot(history.loss, label='Training Set Loss')
plt.title('Epochs vs Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

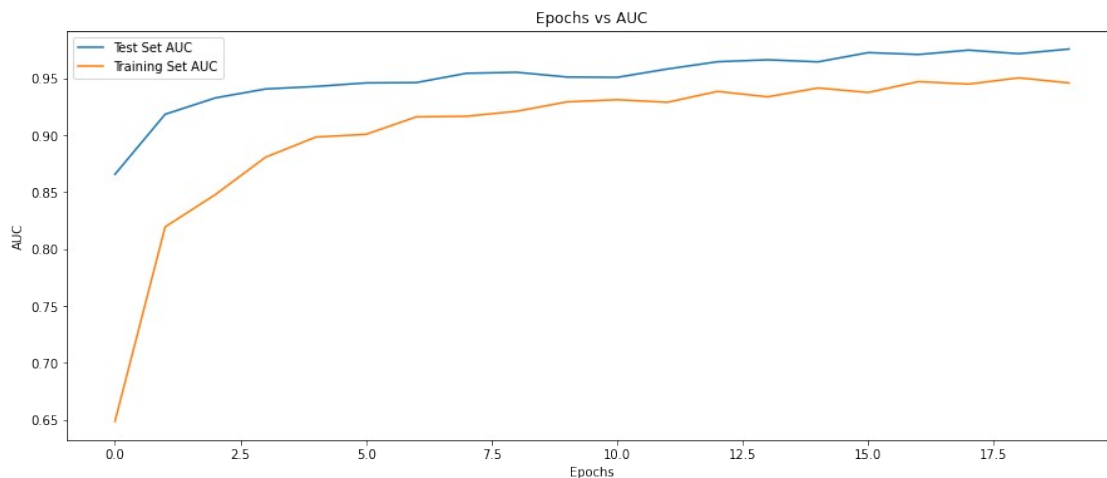
```



```

plt.figure(figsize = (15,6))
plt.plot(history.val_auc, label='Test Set AUC')
plt.plot(history.auc, label='Training Set AUC')
plt.title('Epochs vs AUC')
plt.xlabel('Epochs')
plt.ylabel('AUC')
plt.legend()
plt.show()

```



```

y_train_pred = model.predict(x_train)
cm1 = confusion_matrix(y_train, y_train_pred>0.5)
cm1

```

```

array([[ 847,  343],
       [  59, 5013]])

```

```

y_test_pred = model.predict(x_test)
cm2 = confusion_matrix(y_test, y_test_pred>0.5)
cm2

```

```

array([[ 256,  124],
       [  20, 1688]])

```

```

fpr_train, tpr_train, thresholds_train = roc_curve(y_train,
y_train_pred)
auc_score_train = auc(fpr_train, tpr_train)
print(auc_score_train)

```

```

0.9805026111390929

```

```

fpr_test, tpr_test, thresholds_test = roc_curve(y_test, y_test_pred)
auc_score_test = auc(fpr_test, tpr_test)
print(auc_score_test)

```

```

0.9757780722297548

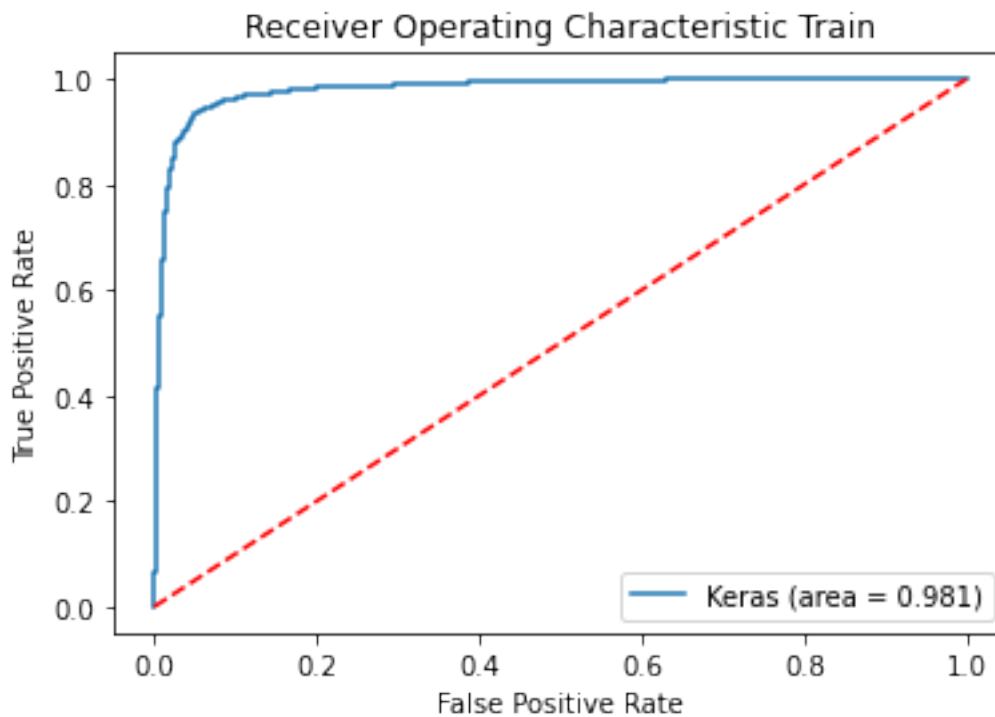
```

```
_, train_acc = model.evaluate(x_train, y_train, verbose=0)
_, test_acc = model.evaluate(x_test, y_test, verbose=0)
print('Train: {}%, Test: {} %'.format((train_acc*100), (test_acc*100)))
```

Train:98.04953336715698%, Test: 97.56362438201904 %


```
# Plot TPR FPR train
```

```
plt.title('Receiver Operating Characteristic Train')
plt.plot(fpr_train, tpr_train, label='Keras (area = {:.3f})'.format(auc_score_train))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
# Plot TPR FPR test
```

```
plt.title('Receiver Operating Characteristic Train')
plt.plot(fpr_test, tpr_test, label='Keras (area = {:.3f})'.format(auc_score_test))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



We received good scores with this model and given data.